

PHASE 2: E&P ASSESSMENT MULTI-AGENT INTEGRATION

Epic-Level Project Plan

PROJECT OVERVIEW

Project Name: E&P Assessment Suggestibility Adapter Agent

Phase: 2 of 4

Platform Integration: HypnoSelfImprovementPlatform + HypnoResilientMind

Ultimate Vision: Integrated Holistic Mental Wellness XR/VR Hub

EPIC BREAKDOWN

EPIC 1: Suggestibility Adapter Agent Foundation

Goal: Create intelligent agent that adapts communication based on E&P Assessment results

Dependencies: Phase 1 E&P Assessment API (COMPLETE)

Timeline: 2-3 weeks

Priority: CRITICAL

EPIC 2: Multi-Agent Orchestration Integration

Goal: Integrate adapter agent with AutoGen/CrewAI multi-agent systems

Dependencies: Epic 1

Timeline: 2 weeks

Priority: HIGH

EPIC 3: Real-Time Language Pattern Adaptation

Goal: Dynamic script modification based on suggestibility type

Dependencies: Epic 1, Epic 2

Timeline: 1-2 weeks

Priority: HIGH

EPIC 4: Quality Assurance & Clinical Safety

Goal: Ensure safe, effective, and clinically appropriate adaptation

Dependencies: Epic 1-3

Timeline: 1 week

Priority: CRITICAL

ARCHITECTURE LAYERS



PRESENTATION LAYER

(React WebXR, Unity/Unreal VR, Web Dashboard)

REST API + WebSocket

API GATEWAY LAYER

FastAPI Backend with E&P Assessment Routes



AI ORCHESTRATION LAYER (PHASE 2 FOCUS)

MULTI-AGENT SYSTEM

- Agent 1: Shadow Recognition Agent
- Agent 2: Clinical Safety Agent
- Agent 3: Script Generation Agent
- Agent 4: Suggestibility Adapter Agent ★ NEW
- Agent 5: Parts Integration Agent
- Agent 6: RAG Query Agent
- Agent 7: Session Orchestrator

COMMUNICATION ADAPTATION ENGINE ★ NEW

- Language Pattern Adapter
- Induction Style Selector
- Script Transformer
- Real-Time Tone Adjuster

DATA LAYER

PostgreSQL + Neo4j + Redis + TimescaleDB + pgvector

EPIC 1: SUGGESTIBILITY ADAPTER AGENT FOUNDATION

Feature 1.1: E&P Preferences Client

Description: Service client to fetch user's communication preferences from E&P API

Story 1.1.1: Create API Client

As a Multi-Agent System

I want to fetch user E&P preferences

So that agents can adapt communication style

Acceptance Criteria:

- Client can call GET /api/v1/assessment/ep/communication-preferences
- Handles authentication (JWT)
- Caches preferences (Redis) for 1 hour
- Handles missing assessments gracefully

Tasks:

- Task 1.1.1.1: Create EPPreferencesClient class
- Task 1.1.1.2: Implement HTTP client with retry logic
- Task 1.1.1.3: Add Redis caching layer
- Task 1.1.1.4: Write unit tests
- Task 1.1.1.5: Create integration tests

Estimated: 1 day

Story 1.1.2: Preferences Response Parser

As a Adapter Agent

I want to parse E&P preferences into structured data

So that I can make adaptation decisions

Acceptance Criteria:

- Parses JSON response from E&P API

- Extracts suggestibility_type, language_patterns, preferred inductions
- Validates data completeness
- Returns typed Python dataclass

Tasks:

- Task 1.1.2.1: Define PreferencesData dataclass
- Task 1.1.2.2: Implement JSON parser
- Task 1.1.2.3: Add validation logic
- Task 1.1.2.4: Write parser tests

Estimated: 0.5 days

Feature 1.2: Language Pattern Adapter

Description: Core engine that transforms generic text to suggestibility-specific language

Story 1.2.1: Physical Suggestible Adapter

As a Script Generator

I want to convert text to Physical-appropriate language

So that Physical suggestibles receive literal, direct communication

Acceptance Criteria:

- Converts inferential to direct language
- Removes metaphors, replaces with literal descriptions
- Emphasizes body sensations
- Uses present-tense, authoritative tone

Example:

Input: "You might find yourself relaxing..."

Output: "Close your eyes. Relax now."

Input: "Like a feather floating down..."

Output: "Feel your body becoming heavy. Notice the weight."

Tasks:

- Task 1.2.1.1: Define Physical transformation rules
- Task 1.2.1.2: Implement pattern matching & replacement
- Task 1.2.1.3: Create transformation pipeline

Task 1.2.1.4: Write transformation tests (20+ examples)

Task 1.2.1.5: Benchmark performance

Estimated: 2 days

Story 1.2.2: Emotional Suggestible Adapter

As a Script Generator

I want to convert text to Emotional-appropriate language

So that Emotional suggestibles receive indirect, metaphorical communication

Acceptance Criteria:

- Converts direct to inferential language
- Adds metaphorical imagery
- Emphasizes emotions and feelings
- Uses past/future tense, permissive tone

Example:

Input: "Close your eyes. Relax now."

Output: "You might find your eyelids becoming heavy... almost as if they prefer to close..."

Input: "Feel your body becoming heavy."

Output: "Perhaps you'll discover a sense of gentle heaviness... like a warm blanket..."

Tasks:

Task 1.2.2.1: Define Emotional transformation rules

Task 1.2.2.2: Implement metaphor generation

Task 1.2.2.3: Create transformation pipeline

Task 1.2.2.4: Write transformation tests (20+ examples)

Task 1.2.2.5: Benchmark performance

Estimated: 2 days

Story 1.2.3: Somnambulistic (Balanced) Adapter

As a Script Generator

I want to use balanced approach

So that Somnambulistic individuals receive appropriate mixed communication

Acceptance Criteria:

- Mixes direct and inferential language

- Balances literal and metaphorical
- Alternates between body and emotion focus
- Uses flexible tense and tone

Tasks:

- Task 1.2.3.1: Define balanced rules
- Task 1.2.3.2: Implement mixed-approach transformer
- Task 1.2.3.3: Write tests

Estimated: 1 day

Feature 1.3: Induction Style Selector

Description: Selects appropriate hypnotic induction based on suggestibility type

Story 1.3.1: Induction Template Library

As a Adapter Agent

I want to have pre-defined induction templates for each type

So that I can select clinically-appropriate techniques

Acceptance Criteria:

- Template library with 3+ inductions per type
- Physical: Progressive Relaxation, Eye Fixation, Body Scan
- Emotional: Guided Imagery, Metaphorical Journey, Confusion
- Somnambulistic: Dave Elman, Rapid, Combined

Tasks:

- Task 1.3.1.1: Create InductionTemplate dataclass
- Task 1.3.1.2: Write 9+ induction templates (HMI-based)
- Task 1.3.1.3: Validate with clinical hypnotherapist
- Task 1.3.1.4: Store in database
- Task 1.3.1.5: Create retrieval service

Estimated: 2 days

Story 1.3.2: Dynamic Induction Selector

As a Adapter Agent

I want to select best induction based on user profile

So that each session uses optimal technique

Acceptance Criteria:

- Considers: suggestibility type, confidence score, past responses
- Returns ranked list of suitable inductions
- Logs selection reasoning

Tasks:

- Task 1.3.2.1: Implement selection algorithm
- Task 1.3.2.2: Add confidence-based filtering
- Task 1.3.2.3: Create ranking system
- Task 1.3.2.4: Write selector tests

Estimated: 1 day

Feature 1.4: Adapter Agent Core

Description: Main agent class that orchestrates E&P-based adaptation

Story 1.4.1: Create SuggestibilityAdapterAgent Class

As a Multi-Agent Orchestrator

I want to instantiate adapter agent

So that it can participate in multi-agent workflows

Acceptance Criteria:

- Compatible with AutoGen AssistantAgent interface
- Compatible with CrewAI Agent interface
- Has tools: adapt_script, select_induction, get_preferences
- Maintains conversation context
- Logs all adaptations

Tasks:

- Task 1.4.1.1: Define agent interface
- Task 1.4.1.2: Implement AutoGen version
- Task 1.4.1.3: Implement CrewAI version
- Task 1.4.1.4: Create tool wrappers
- Task 1.4.1.5: Add logging/monitoring
- Task 1.4.1.6: Write integration tests

Estimated: 3 days

EPIC 2: MULTI-AGENT ORCHESTRATION INTEGRATION

Feature 2.1: AutoGen Integration (ResilientMind Platform)

Description: Integrate adapter agent with existing AutoGen multi-agent system

Story 2.1.1: Add Adapter to Hypno Team

As a ResilientMind Orchestrator

I want to include Suggestibility Adapter in agent team

So that all scripts are personalized

Acceptance Criteria:

- Adapter agent joins AutoGen RoundRobinGroupChat
- Receives user_id at session start
- Fetches E&P preferences before script generation
- Passes preferences to Script Generation Agent

Tasks:

- Task 2.1.1.1: Modify build_hypno_team() to include adapter
- Task 2.1.1.2: Update team workflow (adapter runs first)
- Task 2.1.1.3: Pass preferences via message context
- Task 2.1.1.4: Test full workflow
- Task 2.1.1.5: Verify tool calls in logs

Estimated: 2 days

Story 2.1.2: RAG Integration with E&P

As a RAG Query Agent

I want to filter results by suggestibility type

So that retrieved content is appropriate

Acceptance Criteria:

- RAG queries tagged with suggestibility_type
- Database stores content suitability metadata
- Search results filtered/ranked by appropriateness

Tasks:

- Task 2.1.2.1: Add suggestibility_tags to hypno_documents table
- Task 2.1.2.2: Update RAG query to accept filter parameter
- Task 2.1.2.3: Implement filtered search

Task 2.1.2.4: Test with different types

Estimated: 1.5 days

Feature 2.2: CrewAI Integration

Description: Integrate adapter agent with CrewAI orchestration

Story 2.2.1: Create CrewAI Adapter Agent

As a CrewAI Orchestrator

I want to use adapter as CrewAI agent

So that I can leverage Crew features

Acceptance Criteria:

- Adapter implements CrewAI Agent interface
- Has defined role, goal, backstory
- Tools work with CrewAI task system
- Integrates with other crew members

Tasks:

Task 2.2.1.1: Create CrewAI wrapper class

Task 2.2.1.2: Define agent role/goal/backstory

Task 2.2.1.3: Implement CrewAI tools

Task 2.2.1.4: Write integration tests

Estimated: 2 days

Story 2.2.2: Session Generation Crew with Adapter

As a Therapeutic Session Orchestrator

I want to multi-agent crew with adapter included

So that sessions are personalized and clinically sound

Acceptance Criteria:

- Crew includes: Safety Agent, Adapter Agent, Script Generator
- Sequential workflow: Safety → Adapter → Generator
- Adapter findings passed to downstream agents

Tasks:

Task 2.2.2.1: Define crew structure

Task 2.2.2.2: Create sequential workflow

Task 2.2.2.3: Implement task handoffs

Task 2.2.2.4: Test full crew execution

Estimated: 2 days

Feature 2.3: Session Orchestrator Enhancement

Description: Update session orchestrator to use adapter

Story 2.3.1: Pre-Session E&P Check

As a Session Orchestrator

I want to verify E&P assessment completion before session

So that personalization is always available

Acceptance Criteria:

- Checks if user has completed E&P assessment
- Prompts completion if missing
- Loads preferences before agent initialization
- Caches preferences for session duration

Tasks:

Task 2.3.1.1: Add E&P check to session startup

Task 2.3.1.2: Implement completion prompt UI

Task 2.3.1.3: Cache preferences in session state

Task 2.3.1.4: Add fallback for missing assessment

Estimated: 1.5 days

EPIC 3: REAL-TIME LANGUAGE PATTERN ADAPTATION

Feature 3.1: Script Transformation Pipeline

Description: End-to-end pipeline for transforming generic scripts to personalized

Story 3.1.1: Multi-Stage Transformer

As a Script Generator

I want to transform script through multiple stages

So that adaptation is comprehensive and clinically appropriate

Stages:

1. Parse script structure (intro, induction, deepening, suggestions, emergence)

2. Apply language pattern transformation
3. Adjust pacing and pauses
4. Verify clinical safety
5. Format output

Tasks:

- Task 3.1.1.1: Define pipeline stages
- Task 3.1.1.2: Implement stage interfaces
- Task 3.1.1.3: Create pipeline orchestrator
- Task 3.1.1.4: Add stage validation
- Task 3.1.1.5: Write pipeline tests

Estimated: 2 days

Feature 3.2: Confidence-Based Adaptation

Description: Adjust adaptation strength based on E&P confidence score

Story 3.2.1: Adaptive Transformation Strength

As a Adapter Agent

I want to use lighter adaptation for low-confidence assessments

So that uncertain data doesn't over-influence output

Rules:

- Confidence $\geq 80\%$: Full transformation
- Confidence 60-79%: Moderate transformation
- Confidence $< 60\%$: Minimal transformation + flag for review

Tasks:

- Task 3.2.1.1: Implement confidence-based scaling
- Task 3.2.1.2: Define transformation strength levels
- Task 3.2.1.3: Add clinical review flagging
- Task 3.2.1.4: Test with various confidence scores

Estimated: 1 day

EPIC 4: QUALITY ASSURANCE & CLINICAL SAFETY

Feature 4.1: Clinical Safety Validation

Description: Ensure all adapted content is clinically appropriate

Story 4.1.1: Contraindication Checker

As a Clinical Safety Agent

I want to validate adapted scripts against contraindications

So that harmful content is never produced

Acceptance Criteria:

- Checks for: re-traumatization risks, inappropriate suggestions, unsafe pacing
- Validates induction appropriateness
- Flags for clinician review if needed

Tasks:

- Task 4.1.1.1: Define contraindication rules
- Task 4.1.1.2: Implement checker logic
- Task 4.1.1.3: Integrate with adapter workflow
- Task 4.1.1.4: Create safety tests

Estimated: 2 days

Feature 4.2: Adaptation Quality Metrics

Description: Track and monitor adaptation effectiveness

Story 4.2.1: Metrics Collection

As a Platform Administrator

I want to track adaptation quality metrics

So that I can improve system performance

Metrics:

- Adaptation rate (% sessions using adapter)
- Confidence distribution
- User satisfaction by suggestibility type
- Clinical review rate

Tasks:

- Task 4.2.1.1: Define metrics schema

- Task 4.2.1.2: Implement collection points
- Task 4.2.1.3: Create analytics dashboard
- Task 4.2.1.4: Set up alerting

Estimated: 2 days

TIMELINE SUMMARY

Epic	Duration	Dependencies
Epic 1: Foundation	2-3 weeks	Phase 1 complete
Epic 2: Integration	2 weeks	Epic 1
Epic 3: Adaptation	1-2 weeks	Epic 1, 2
Epic 4: Safety/QA	1 week	Epic 1-3
TOTAL	6-8 weeks	Sequential

INTEGRATION WITH RESILIENTMIND PLATFORM

Current ResilientMind Architecture:

User Question → FastAPI Core → AutoGen Agent → MCP RAG Server → PostgreSQL

With Phase 2 Integration:

User Question → FastAPI Core →

↓

1. E&P Preferences Client → E&P Assessment API

↓

2. AutoGen Team (with Adapter Agent) →

- Safety Agent (checks contraindications)
- Adapter Agent (fetches E&P, adapts approach) ★ NEW
- RAG Agent (retrieves filtered content)
- Script Generator (produces personalized script)

↓

3. MCP RAG Server → PostgreSQL (filtered by suggestibility)

↓

4. Response (personalized, adapted, safe)

SUCCESS CRITERIA

Phase 2 is successful when:

- Adapter agent operational in both AutoGen and CrewAI
 - 90%+ of sessions use E&P-based personalization
 - Language patterns appropriately adapted (validated by clinicians)
 - No adverse clinical safety events
 - User satisfaction increases by 25%+
 - Integration tests pass for both platforms
-

NEXT STEPS

Immediate Actions:

1. Review this plan with clinical team
2. Set up project management (Jira/Linear)
3. Create development branch
4. Begin Epic 1, Feature 1.1 (E&P Preferences Client)

Ready to start Phase 2 implementation?